



1 Representation learning



$$1024-D = 256 + 256 + 512$$

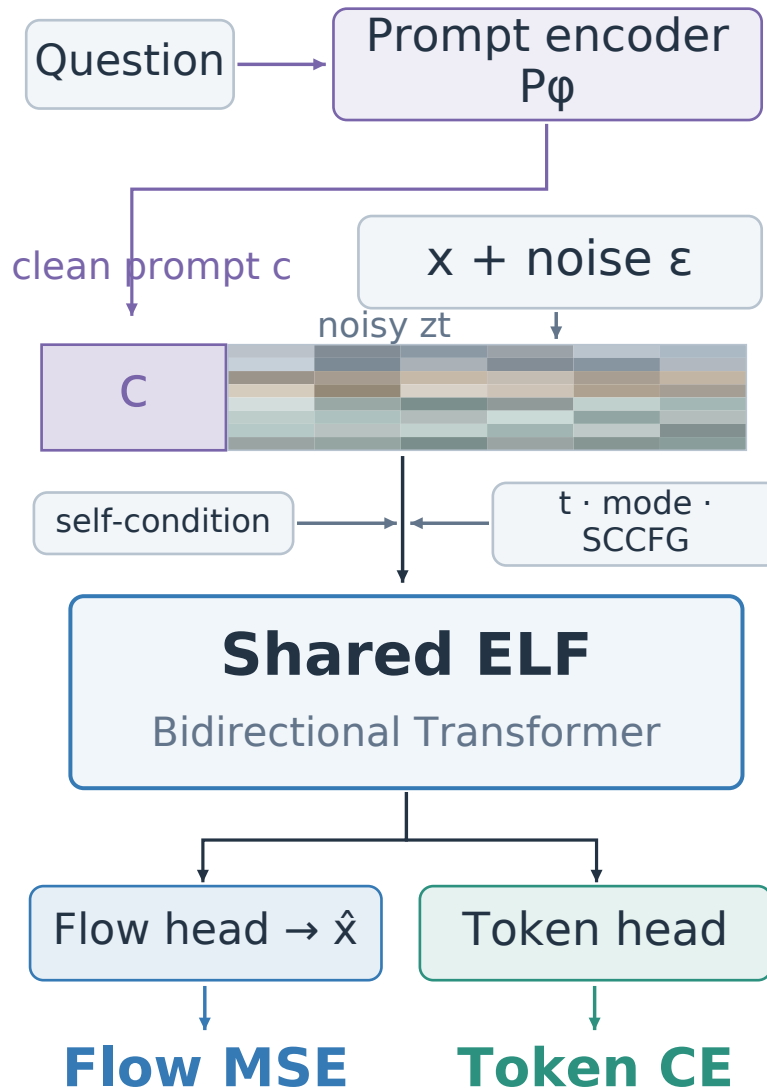
$$x_l = (h_l - \mu_l) E_l$$

Teacher-CE supervision



Covariance whitening; then freeze the representation for flow training.

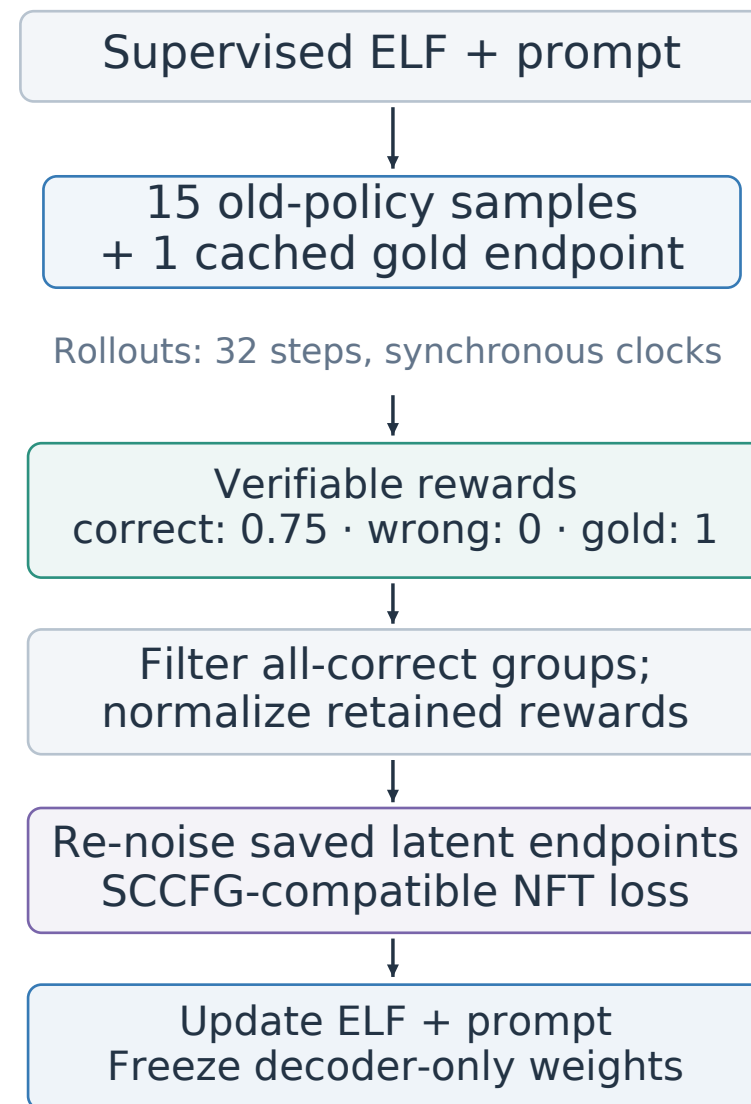
2 Flow training



80% flow rows · 20% decoder rows

Teacher prompt → MSE warm start → joint

3 NFT post-training



Forward-process training;
fixed reference-policy penalty