



1 Representation learning



$$1024-D = 256 + 256 + 512$$


$$x_l = (h_l - \mu_l) E_l$$

Teacher-CE supervision





Covariance whitening; then freeze the representation for flow training.

2 Flow training



3 Async inference

Encode question once → cache c



Joint ELF calls + Euler updates

ELF + token head → answer

$$dz_l/d\tau = f_l(\tau) v_l \cdot \text{decode once}$$