


Measured inference cost versus total denoising steps

8 steps/sample 16 steps/sample 32 steps/sample 64 steps/sample



k × measured amortized GPU-seconds/answer; includes prompt, denoising and decoding; excludes startup, scoring and voting.
One H200, batch32 per case. ELF-B NFT: two concurrent cases, overlap counted once. Other panels: one case/H200. Not single-request latency.